

## >Notebook 3: Nominal Models Final Project - Ordinal vs Nominal Sentiment Analysis

Atharv Chaudhary

**Purpose:** Train and evaluate NOMINAL classification models.

### Models:

1. Multinomial Naive Bayes
2. Logistic Regression (Multinomial)

**Input:** `amazon_electronics_cleaned.csv`

**Output:** `nominal_results.csv`, confusion matrices

```
1 from google.colab import drive
2 drive.mount('/content/drive')

Drive already mounted at /content/drive; to attempt to forcibly remount, call dr
```

```
1 # Import libraries
2 import pandas as pd
3 import numpy as np
4 import matplotlib.pyplot as plt
5 import seaborn as sns
6 import warnings
7 warnings.filterwarnings('ignore')
8
9 from sklearn.model_selection import train_test_split
10 from sklearn.feature_extraction.text import TfidfVectorizer
11 from sklearn.naive_bayes import MultinomialNB
12 from sklearn.linear_model import LogisticRegression
13 from sklearn.metrics import (
14     accuracy_score,
15     mean_absolute_error,
16     f1_score,
17     classification_report,
18     confusion_matrix
19 )
20
21 # Settings
22 RANDOM_STATE = 42
23 np.random.seed(RANDOM_STATE)
24 plt.style.use('seaborn-v0_8-whitegrid')
25
26 print("✅ Libraries imported")
```

✅ Libraries imported

## Step 1: Load Data

```
1 # Load cleaned data
2 df = pd.read_csv('/content/drive/MyDrive/Fall 2025/Foundations of Artificial Intelligence/reviews.csv')
3 print(f"✅ Loaded {len(df)} reviews")
4
5 # Show distribution
6 print("\n📊 Rating Distribution:")
7 print(df['rating'].value_counts().sort_index())
```

✅ Loaded 49,960 reviews

📊 Rating Distribution:

| rating | count |
|--------|-------|
| 1      | 2835  |
| 2      | 2161  |
| 3      | 3964  |
| 4      | 10103 |
| 5      | 30897 |

Name: count, dtype: int64

## Step 2: Feature Extraction (TF-IDF)

```
1 # =====
2 # TF-IDF FEATURE EXTRACTION
3 # =====
4
5 print("=" * 70)
6 print("TF-IDF FEATURE EXTRACTION")
7 print("=" * 70)
8
9 # Configuration
10 MAX_FEATURES = 5000
11 NGRAM_RANGE = (1, 2) # Unigrams and bigrams
12
13 vectorizer = TfidfVectorizer(
14     max_features=MAX_FEATURES,
15     stop_words='english',
16     ngram_range=NGRAM_RANGE,
17     min_df=5,
18     max_df=0.95
19 )
20
21 print(f"\nSettings:")
22 print(f"  Max features: {MAX_FEATURES}")
23 print(f"  N-gram range: {NGRAM_RANGE}")
24
```

```

25 # Transform
26 X = vectorizer.fit_transform(df['text'])
27 y = df['rating'].values
28
29 print(f"\n Training set: {X_train.shape[0]:,} samples")
13 print(f"<span style='color: green; font-size: 1.5em; vertical-align: middle; margin-right: 0.1em; border: 1px solid green; padding: 0 2px; font-weight: bold; font-style: italic; font-family: inherit; line-height: 1; margin-right: 0.1em; border-radius: 50%; width: 1em; height: 1em; display: inline-block;"> Test set: {X_test.shape[0]:,} samples")

```

✓ Training set: 39,968 samples

✓ Test set: 9,992 samples

## Step 4: Helper Functions

```

1 # =====
2 # HELPER FUNCTIONS
3 # =====
4
5 def evaluate_model(y_true, y_pred, model_name):

```

```

6     """Evaluate model and return metrics."""
7     accuracy = accuracy_score(y_true, y_pred)
8     mae = mean_absolute_error(y_true, y_pred)
9     f1_macro = f1_score(y_true, y_pred, average='macro')
10    f1_weighted = f1_score(y_true, y_pred, average='weighted')
11
12    print(f"\n{'='*55}")
13    print(f"📊 {model_name}")
14    print(f"{'='*55}")
15    print(f"Accuracy:      {accuracy:.4f} ({accuracy*100:.2f}%)")
16    print(f"MAE:           {mae:.4f}")
17    print(f"F1 (macro):   {f1_macro:.4f}")
18    print(f"F1 (weighted): {f1_weighted:.4f}")
19
20    return {
21        'model': model_name,
22        'encoding': 'Nominal',
23        'accuracy': accuracy,
24        'mae': mae,
25        'f1_macro': f1_macro,
26        'f1_weighted': f1_weighted
27    }
28
29
30 def calculate_error_rates(y_true, y_pred):
31     """Calculate adjacent and severe error rates."""
32     errors = y_true != y_pred
33     if errors.sum() == 0:
34         return 0.0, 0.0
35
36     error_distances = np.abs(y_true[errors] - y_pred[errors])
37     adjacent = (error_distances == 1).sum() / errors.sum()
38     severe = (error_distances >= 2).sum() / errors.sum()
39
40     return adjacent, severe
41
42
43 print("✅ Helper functions defined")

```

Helper functions defined

## Step 5: Model 1 - Multinomial Naive Bayes

```

1 # =====
2 # MODEL 1: MULTINOMIAL NAIVE BAYES
3 # =====

```

```

4
5 print("\n" + "="*70)
6 print("🔧 MODEL 1: Multinomial Naive Bayes (Nominal)")
7 print("=*70)
8 print("\nTreats classes as UNORDERED categories.")
9 print("Formula:  $P(Y=k|x) \propto P(Y=k) \times \prod P(x_j|Y=k)$ ")
10
11 # Train
12 nb_model = MultinomialNB(alpha=1.0) # Laplace smoothing
13 nb_model.fit(X_train, y_train)
14
15 # Predict
16 nb_pred = nb_model.predict(X_test)
17
18 # Evaluate
19 nb_results = evaluate_model(y_test, nb_pred, "Naive Bayes")
20
21 # Error analysis
22 nb_adjacent, nb_severe = calculate_error_rates(y_test, nb_pred)
23 nb_results['adjacent_error'] = nb_adjacent
24 nb_results['severe_error'] = nb_severe
25
26 print(f"\nError Analysis:")
27 print(f"    Adjacent Error Rate ( $\pm 1$ ): {nb_adjacent:.2%}")

```

```

=====
🔧 MODEL 1: Multinomial Naive Bayes (Nominal)
=====

Treats classes as UNORDERED categories.
Formula:  $P(Y=k|x) \propto P(Y=k) \times \prod P(x_j|Y=k)$ 

=====
📊 Naive Bayes
=====

Accuracy: 0.6312 (63.12%)
MAE: 0.6651
F1 (macro): 0.2321
F1 (weighted): 0.5129

Error Analysis:
    Adjacent Error Rate ( $\pm 1$ ): 55.63%
    Severe Error Rate ( $\pm 2+$ ): 44.37%

```

```

1 # Classification report
2 print("\n📋 Classification Report - Naive Bayes:")
3 print(classification_report(y_test, nb_pred, digits=4))

```

| Classification Report - Naive Bayes: |           |        |          |         |
|--------------------------------------|-----------|--------|----------|---------|
|                                      | precision | recall | f1-score | support |
| 1                                    | 0.6236    | 0.1958 | 0.2980   | 567     |
| 2                                    | 0.0000    | 0.0000 | 0.0000   | 432     |
| 3                                    | 0.3750    | 0.0038 | 0.0075   | 793     |
| 4                                    | 0.3797    | 0.0445 | 0.0797   | 2021    |
| 5                                    | 0.6378    | 0.9877 | 0.7751   | 6179    |
| accuracy                             |           |        | 0.6312   | 9992    |
| macro avg                            | 0.4032    | 0.2464 | 0.2321   | 9992    |
| weighted avg                         | 0.5364    | 0.6312 | 0.5129   | 9992    |

## Step 6: Model 2 - Logistic Regression (Multinomial)

```

1 # =====
2 # MODEL 2: LOGISTIC REGRESSION (MULTINOMIAL)
3 # =====
4
5 print("\n" + "="*70)
6 print("🔧 MODEL 2: Logistic Regression (Nominal - Multinomial)")
7 print("="*70)
8 print("\nUses softmax, treats classes as UNORDERED.")
9 print("Formula:  $P(Y=k|x) = \exp(w_k^T x + b_k) / \sum \exp(w_j^T x + b_j)$ ")
10
11 # Train
12 lr_model = LogisticRegression(
13     multi_class='multinomial',
14     solver='lbfgs',
15     max_iter=1000,
16     random_state=RANDOM_STATE,
17     n_jobs=-1
18 )
19 lr_model.fit(X_train, y_train)
20
21 # Predict
22 lr_pred = lr_model.predict(X_test)
23
24 # Evaluate
25 lr_results = evaluate_model(y_test, lr_pred, "Logistic Regression")
26
27 # Error analysis
28 lr_adjacent, lr_severe = calculate_error_rates(y_test, lr_pred)
29 lr_results['adjacent_error'] = lr_adjacent
30 lr_results['severe_error'] = lr_severe
31
32 print(f"\nError Analysis:")
33 print(f"  Adjacent Error Rate (\pm1): {lr_adjacent:.2%}")
34 print(f"  Severe Error Rate (\pm2+): {lr_severe:.2%}")

```

=====

🔧 MODEL 2: Logistic Regression (Nominal - Multinomial)

=====

Uses softmax, treats classes as UNORDERED.

Formula:  $P(Y=k|x) = \exp(w_k^T x + b_k) / \sum \exp(w_j^T x + b_j)$

=====

 Logistic Regression

=====

Accuracy: 0.6595 (65.95%)

MAE: 0.5337

```
F1 (macro): 0.3793
F1 (weighted): 0.6064
```

## Error Analysis:

Adjacent Error Rate ( $\pm 1$ ): 65.17%  
Severe Error Rate ( $\pm 2+$ ): 34.83%

```
1 # Classification report
2 print("\n📋 Classification Report - Logistic Regression:")
3 print(classification_report(y_test, lr_pred, digits=4))
```

| 📋 Classification Report - Logistic Regression: |           |        |          |         |
|--|-----------|--------|----------|---------|
|  | precision | recall | f1-score | support |
| 1  | 0.6108    | 0.4374 | 0.5098   | 567     |
| 2  | 0.3059    | 0.0602 | 0.1006   | 432     |
| 3  | 0.3404    | 0.1223 | 0.1800   | 793     |
| 4  | 0.4197    | 0.2340 | 0.3005   | 2021    |
| 5  | 0.7103    | 0.9299 | 0.8054   | 6179    |
| accuracy                                       |           |        | 0.6595   | 9992    |
| macro avg                                      | 0.4774    | 0.3568 | 0.3793   | 9992    |
| weighted avg                                   | 0.5991    | 0.6595 | 0.6064   | 9992    |

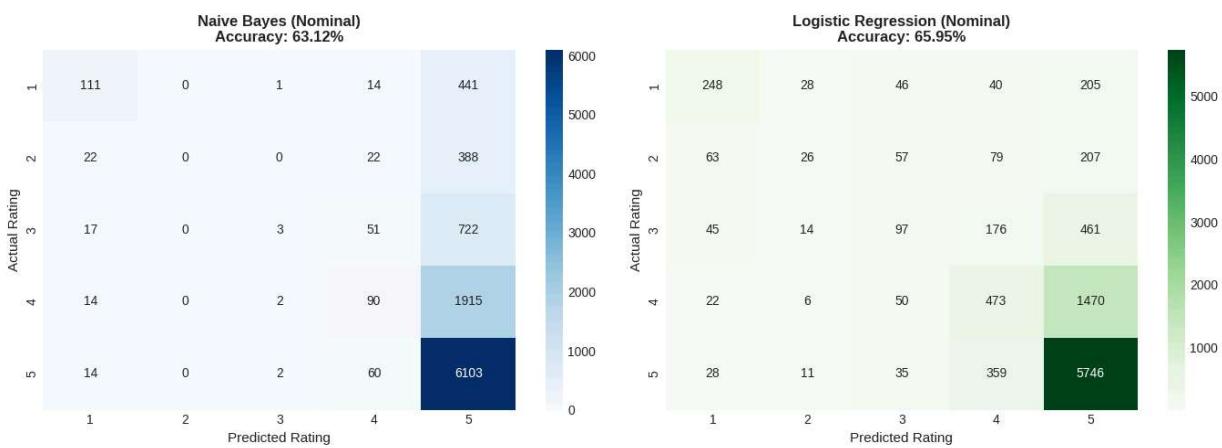
## Step 7: Confusion Matrices

```
1 # =====
2 # CONFUSION MATRICES
3 # =====
4
5 fig, axes = plt.subplots(1, 2, figsize=(14, 5))
6
7 # Naive Bayes
8 cm_nb = confusion_matrix(y_test, nb_pred)
9 sns.heatmap(cm_nb, annot=True, fmt='d', cmap='Blues', ax=axes[0],
10             xlabel='Predicted Rating', fontsize=11)
11 axes[0].set_ylabel('Actual Rating', fontsize=11)
12 axes[0].set_title(f'Naive Bayes (Nominal)\nAccuracy: {nb_results["accuracy"]}', 
13                   fontsize=12, fontweight='bold')
14
15 # Logistic Regression
16 cm_lr = confusion_matrix(y_test, lr_pred)
17 sns.heatmap(cm_lr, annot=True, fmt='d', cmap='Greens', ax=axes[1],
18             xlabel='Predicted Rating', fontsize=11)
19 axes[1].set_ylabel('Actual Rating', fontsize=11)
20 axes[1].set_title(f'Logistic Regression (Nominal)\nAccuracy: {lr_results["accuracy"]}', 
21                   fontsize=12, fontweight='bold')
```

```

24
25 plt.tight_layout()
26 plt.savefig('confusion_matrices_nominal.png', dpi=150, bbox_inches='tight'
27 plt.show()
28
29 print("\n✅ Saved: confusion_matrices_nominal.png")

```



Saved: confusion\_matrices\_nominal.png

## Step 8: Save Results

```

1 # =====
2 # SAVE RESULTS
3 # =====
4
5 # Combine results
6 nominal_results = pd.DataFrame([nb_results, lr_results])
7
8 print("\n" + "="*70)
9 print("📊 NOMINAL MODELS SUMMARY")
10 print("="*70)
11 print(nominal_results.to_string(index=False))
12
13 # Save to CSV
14 nominal_results.to_csv('nominal_results.csv', index=False)
15 print("\n✅ Saved: nominal_results.csv")

```

```
=====
 NOMINAL MODELS SUMMARY
=====
model encoding accuracy mae f1_macro f1_weighted adjacent
Naive Bayes Nominal 0.631205 0.665132 0.232055 0.512934 0.
Logistic Regression Nominal 0.659528 0.533727 0.379251 0.606418 0.
```

Saved: nominal\_results.csv

```
1 # Save predictions for later analysis
2 predictions_df = pd.DataFrame({
3     'actual': y_test,
4     'nb_pred': nb_pred,
5     'lr_pred': lr_pred
6 })
7 predictions_df.to_csv('nominal_predictions.csv', index=False)
8 print(" Saved: nominal_predictions.csv")
```

Saved: nominal\_predictions.csv

```
1 # Download files
2 try:
3     from google.colab import files
4     files.download('nominal_results.csv')
5     files.download('confusion_matrices_nominal.png')
6 except:
7     print("Files saved locally")
```

```
1 import shutil
2 import os
3
4 drive_path = '/content/drive/MyDrive/Fall 2025/Foundations of Artificial
Intelligence/Final Project/data'
5
6 # Create the directory if it doesn't exist
7 os.makedirs(drive_path, exist_ok=True)
8
9 files_to_save = [
10     'nominal_results.csv',
11     'confusion_matrices_nominal.png',
12     'nominal_predictions.csv'
13 ]
14
15 for file_name in files_to_save:
16     source_path = os.path.join('/content', file_name)
17     destination_path = os.path.join(drive_path, file_name)
18     try:
```

```
19     shutil.copy(source_path, destination_path)
20     print(f"✅ Saved '{file_name}' to Drive: {destination_path}")
21 except FileNotFoundError:
22     print(f"⚠️ Error: '{file_name}' not found. Skipping.")
23 except Exception as e:
24     print(f"❌ Error saving '{file_name}' to Drive: {e}")
```

✓ Saved 'nominal\_results.csv' to Drive: /content/drive/MyDrive/Fall 2025/Foundations of Data Science/nominal\_results.csv  
✓ Saved 'confusion\_matrices\_nominal.png' to Drive: /content/drive/MyDrive/Fall 2025/Foundations of Data Science/confusion\_matrices\_nominal.png  
✓ Saved 'nominal\_predictions.csv' to Drive: /content/drive/MyDrive/Fall 2025/Foundations of Data Science/nominal\_predictions.csv

---

## Summary