

```

1 print("Installing required packages...")
2 import sys
3 import subprocess
4
5 packages = ['datasets', 'scikit-learn', 'pandas', 'numpy']
6 for package in packages:
7     try:
8         __import__(package.replace('-', '_'))
9     except ImportError:
10        print(f"Installing {package}...")
11        subprocess.check_call([sys.executable, "-m", "pip", "install", package, "-q"])
12
13 print("✓ All packages installed!")
14

```

Installing required packages...
 Installing scikit-learn...
 ✓ All packages installed!

```

1 import pandas as pd
2 import numpy as np
3 from sklearn.model_selection import train_test_split
4 from sklearn.feature_extraction.text import TfidfVectorizer
5 from sklearn.naive_bayes import MultinomialNB
6 from sklearn.linear_model import LogisticRegression
7 from sklearn.metrics import accuracy_score, confusion_matrix, classification_report, mean_absolute_error
8 import re
9 import warnings
10 warnings.filterwarnings('ignore')
11
12 print("✓ Libraries imported successfully")

```

✓ Libraries imported successfully

```

1 print("=" * 70)
2 print("LOADING AMAZON ELECTRONICS REVIEWS DATASET")
3 print("=" * 70)

```

=====
 LOADING AMAZON ELECTRONICS REVIEWS DATASET
 =====

```

1 df = None
2 try:
3     from datasets import load_dataset
4
5     print("\nLoading Electronics reviews from HuggingFace...")
6     print("This may take 2-3 minutes on first download...\n")
7
8     dataset = load_dataset(
9         "McAuley-Lab/Amazon-Reviews-2023",
10        "raw_review_Electronics",
11        split="full[:10000]",
12        trust_remote_code=False  # Updated parameter
13    )
14
15    df = pd.DataFrame(dataset)
16    print(f"✓ Loaded {len(df)} reviews via HuggingFace")
17
18 except Exception as e:
19     print(f"HuggingFace method failed: {e}")
20     print("\nTrying alternative: Old Amazon Dataset (2014)...\n")
21
22     # Method 2: Old UCSD Dataset (backup)
23     try:
24         import gzip
25         import json
26         import urllib.request
27
28         url = "http://snap.stanford.edu/data/amazon/productGraph/categoryFiles/reviews_Electronics_5.json.gz"
29
30         print("Downloading from UCSD...")
31         urllib.request.urlretrieve(url, 'electronics.json.gz')
32

```

```

33     reviews = []
34     with gzip.open('electronics.json.gz', 'rt', encoding='utf-8') as f:
35         for i, line in enumerate(f):
36             if i >= 10000:
37                 break
38             reviews.append(json.loads(line))
39
40     df = pd.DataFrame(reviews)
41     print(f"✓ Loaded {len(df)} reviews from 2014 dataset")
42
43 except Exception as e2:
44     print(f"✗ All methods failed: {e2}")
45     raise
46
47 # Display dataset info
48 print(f"\nDataset Shape: {df.shape}")
49 print(f"Columns: {df.columns.tolist()}")
50 print("\nFirst 3 reviews:")
51 print(df.head(3))

```

Loading Electronics reviews from HuggingFace...
This may take 2-3 minutes on first download...

README.md: 30.3k? [00:00<00:00, 1.33MB/s]

Amazon-Reviews-2023.py: 39.6k? [00:00<00:00, 1.57MB/s]

HuggingFace method failed: Dataset scripts are no longer supported, but found Amazon-Reviews-2023.py

Trying alternative: Old Amazon Dataset (2014)...

Downloading from UCSD...

✓ Loaded 10000 reviews from 2014 dataset

Dataset Shape: (10000, 9)

Columns: ['reviewerID', 'asin', 'reviewerName', 'helpful', 'reviewText', 'overall', 'summary', 'unixReviewTime', 'reviewTime']

First 3 reviews:

	reviewerID	asin	reviewerName	helpful	reviewText	overall
0	A094DHG771SJ	0528881469	amazdnu	[0, 0]	We got this GPS for my husband who is an (OTR)...	5.0
1	AM0214LNFCIEI4	0528881469	Amazon Customer	[12, 15]	I'm a professional OTR truck driver, and I bou...	1.0
2	A3N7T0DY83Y4IG	0528881469	C. A. Freeman	[43, 45]	Well, what can I say. I've had this unit in m...	3.0

	summary	unixReviewTime	reviewTime
0	Gotta have GPS!	1370131200	06 2, 2013
1	Very Disappointed	1290643200	11 25, 2010
2	1st impression	1283990400	09 9, 2010

```

1 print("\n" + "=" * 70)
2 print("DATA PREPROCESSING")
3 print("=" * 70)
4
5 # Identify column names
6 text_col = None
7 rating_col = None
8
9 for col in ['text', 'reviewText', 'review_text']:
10    if col in df.columns:
11        text_col = col
12        break
13
14 for col in ['rating', 'overall', 'stars']:
15    if col in df.columns:
16        rating_col = col
17        break
18
19 print(f"\n✓ Text column: '{text_col}'")
20 print(f"✓ Rating column: '{rating_col}'")
21
22 # Clean data
23 df_clean = df[[text_col, rating_col]].dropna()
24 print(f"✓ Removed {len(df) - len(df_clean)} rows with missing values")
25
26 # Convert ratings to integers

```

```

27 df_clean[rating_col] = pd.to_numeric(df_clean[rating_col], errors='coerce')
28 df_clean = df_clean.dropna()
29 df_clean[rating_col] = df_clean[rating_col].astype(int)
30 df_clean = df_clean[df_clean[rating_col].isin([1, 2, 3, 4, 5])]
31
32 print(f"✓ Final dataset: {len(df_clean)} reviews")
33
34 # Show class distribution
35 print("\n📊 Class Distribution:")
36 dist = df_clean[rating_col].value_counts().sort_index()
37 for rating, count in dist.items():
38     pct = (count / len(df_clean)) * 100
39     print(f"  {rating}-star: {count} ({pct:.5f}%)")

```

=====
DATA PREPROCESSING
=====

✓ Text column: 'reviewText'
✓ Rating column: 'overall'
✓ Removed 0 rows with missing values
✓ Final dataset: 10000 reviews

📊 Class Distribution:
1-star: 572 (5.7%)
2-star: 450 (4.5%)
3-star: 822 (8.2%)
4-star: 2095 (20.9%)
5-star: 6061 (60.6%)

```

1 # %% Cell 5: Text Cleaning
2 print("\n" + "=" * 70)
3 print("TEXT PREPROCESSING")
4 print("=" * 70)
5
6 def clean_text(text):
7     """Clean and normalize text"""
8     if not isinstance(text, str):
9         return ""
10    text = text.lower()
11    text = re.sub(r'[^\w\s]', '', text)
12    text = ' '.join(text.split())
13    return text
14
15 print("\nCleaning text...")
16 df_clean['cleaned_text'] = df_clean[text_col].apply(clean_text)
17
18 # Remove very short reviews
19 df_clean = df_clean[df_clean['cleaned_text'].str.len() >= 10]
20 print(f"✓ Text cleaned. Final reviews: {len(df_clean)}")
21
22 print("\nExample cleaned review:")
23 print(f"Original: {df_clean[text_col].iloc[0][:100]}...")
24 print(f"Cleaned: {df_clean['cleaned_text'].iloc[0][:100]}...")
25

```

=====
TEXT PREPROCESSING
=====

Cleaning text...
✓ Text cleaned. Final reviews: 9991

Example cleaned review:
Original: We got this GPS for my husband who is an (OTR) over the road trucker. Very Impressed with the shipp...
Cleaned: we got this gps for my husband who is an otr over the road trucker very impressed with the shipping ...

```

1 print("\n" + "=" * 70)
2 print("FEATURE EXTRACTION")
3 print("=" * 70)
4
5 X = df_clean['cleaned_text']
6 y = df_clean[rating_col]
7
8 # Train-test split
9 print("\nSplitting data (80% train, 20% test)...")
10 X_train, X_test, y_train, y_test = train_test_split(

```

```

11     X, y, test_size=0.2, random_state=42, stratify=y
12 )
13
14 # TF-IDF Vectorization
15 print("Extracting TF-IDF features...")
16 vectorizer = TfidfVectorizer(max_features=1000, min_df=5, max_df=0.8)
17 X_train_tfidf = vectorizer.fit_transform(X_train)
18 X_test_tfidf = vectorizer.transform(X_test)
19
20 print(f"✓ Training set: {X_train_tfidf.shape[0]} samples")
21 print(f"✓ Test set: {X_test_tfidf.shape[0]} samples")
22 print(f"✓ Features: {X_train_tfidf.shape[1]} TF-IDF terms")
23
24 print("\nTop 10 features by TF-IDF score:")
25 feature_names = vectorizer.get_feature_names_out

```

```
=====
FEATURE EXTRACTION
=====
```

Splitting data (80% train, 20% test)...
 Extracting TF-IDF features...
 ✓ Training set: 7992 samples
 ✓ Test set: 1999 samples
 ✓ Features: 1000 TF-IDF terms

Top 10 features by TF-IDF score:

```

1 print("\n" + "=" * 70)
2 print("TRAINING MODELS")
3 print("=" * 70)
4
5 # Model 1: Multinomial Naive Bayes
6 print("\n[1/2] Training Naive Bayes...")
7 nb_model = MultinomialNB()
8 nb_model.fit(X_train_tfidf, y_train)
9 nb_pred = nb_model.predict(X_test_tfidf)
10 print("✓ Naive Bayes trained")
11
12 # Model 2: Logistic Regression
13 print("\n[2/2] Training Logistic Regression...")
14 lr_model = LogisticRegression(max_iter=500, random_state=42, multi_class='multinomial')
15 lr_model.fit(X_train_tfidf, y_train)
16 lr_pred = lr_model.predict(X_test_tfidf)
17 print("✓ Logistic Regression trained")

```

```
=====
TRAINING MODELS
=====
```

[1/2] Training Naive Bayes...
 ✓ Naive Bayes trained
 [2/2] Training Logistic Regression...
 ✓ Logistic Regression trained

```

1 # %% Cell 8: Evaluate Results
2 print("\n" + "=" * 70)
3 print("EVALUATION RESULTS")
4 print("=" * 70)
5
6 # Calculate metrics
7 nb_acc = accuracy_score(y_test, nb_pred)
8 nb_mae = mean_absolute_error(y_test, nb_pred)
9 nb_cm = confusion_matrix(y_test, nb_pred)
10
11 lr_acc = accuracy_score(y_test, lr_pred)
12 lr_mae = mean_absolute_error(y_test, lr_pred)
13 lr_cm = confusion_matrix(y_test, lr_pred)
14
15 print(f"\n📊 MULTINOMIAL NAIVE BAYES:")
16 print(f"    Accuracy: {nb_acc*100:.1f}% ({nb_acc:.4f})")
17 print(f"    MAE: {nb_mae:.2f}")
18
19 print(f"\n📊 LOGISTIC REGRESSION (Nominal):")
20 print(f"    Accuracy: {lr_acc*100:.1f}% ({lr_acc:.4f})")

```

```

21 print(f"    MAE: {lr_mae:.2f}")
22
23 # Confusion matrices
24 print("\n📊 Confusion Matrix - Naive Bayes:")
25 print(nb_cm)
26
27 print("\n📊 Confusion Matrix - Logistic Regression:")
28 print(lr_cm)
29

```

```
=====
EVALUATION RESULTS
=====
```

📊 MULTINOMIAL NAIVE BAYES:

Accuracy: 60.5% (0.6053)
MAE: 0.74

📊 LOGISTIC REGRESSION (Nominal):

Accuracy: 62.6% (0.6258)
MAE: 0.62

📊 Confusion Matrix - Naive Bayes:

```
[[ 0   0   0   0  114]
 [ 0   0   0   1  89]
 [ 0   0   0   1 164]
 [ 0   0   0   0  419]
 [ 0   0   0   1 1210]]
```

📊 Confusion Matrix - Logistic Regression:

```
[[ 23   2   5   6   78]
 [  9   3   6  20   52]
 [  6   1   6  47 105]
 [  2   1   3  92 321]
 [  2   0   4  78 1127]]
```

```

1 print("\n" + "=" * 70)
2 print("DETAILED ANALYSIS")
3 print("=" * 70)
4
5 # Adjacent rating error analysis
6 def calc_adjacent_errors(cm):
7     """Calculate % of errors that are adjacent ratings"""
8     total_errors = np.sum(cm) - np.trace(cm)
9     if total_errors == 0:
10         return 0
11     adjacent_errors = 0
12     for i in range(len(cm)):
13         for j in range(len(cm)):
14             if abs(i - j) == 1:
15                 adjacent_errors += cm[i][j]
16     return (adjacent_errors / total_errors) * 100
17
18 nb_adj = calc_adjacent_errors(nb_cm)
19 lr_adj = calc_adjacent_errors(lr_cm)
20
21 print(f"\n📊 Adjacent Rating Errors:")
22 print(f"    Naive Bayes: {nb_adj:.1f}% of errors are adjacent (e.g., 4*5)")
23 print(f"    Logistic Regression: {lr_adj:.1f}% of errors are adjacent")
24
25 # Per-class performance
26 print("\n📊 Per-Class F1-Scores:")
27 nb_report = classification_report(y_test, nb_pred, output_dict=True, zero_division=0)
28 lr_report = classification_report(y_test, lr_pred, output_dict=True, zero_division=0)
29
30 print("\nNaive Bayes:")
31 for rating in [1, 2, 3, 4, 5]:
32     if str(rating) in nb_report:
33         f1 = nb_report[str(rating)]['f1-score']
34         support = nb_report[str(rating)]['support']
35         print(f"    {rating}-star: F1 = {f1:.3f} (n={int(support)}))")
36
37 print("\nLogistic Regression:")
38 for rating in [1, 2, 3, 4, 5]:
39     if str(rating) in lr_report:
40         f1 = lr_report[str(rating)]['f1-score']
41         support = lr_report[str(rating)]['support']

```

DETAILED ANALYSIS

Adjacent Rating Errors:
 Naive Bayes: 53.4% of errors are adjacent (e.g., 4→5)
 Logistic Regression: 62.4% of errors are adjacent

Per-Class F1-Scores:

Naive Bayes:
 1-star: F1 = 0.000 (n=114)
 2-star: F1 = 0.000 (n=90)
 3-star: F1 = 0.000 (n=165)
 4-star: F1 = 0.000 (n=419)
 5-star: F1 = 0.755 (n=1211)

Logistic Regression:
 1-star: F1 = 0.295 (n=114)
 2-star: F1 = 0.062 (n=90)
 3-star: F1 = 0.063 (n=165)
 4-star: F1 = 0.278 (n=419)
 5-star: F1 = 0.779 (n=1211)

```

1 # Get key metrics
2 f1_5star = lr_report.get('5', {}).get('f1-score', 0)
3 f1_3star = lr_report.get('3', {}).get('f1-score', 0)
4
5 paragraph = f"""Initial experiments on {len(df_clean):,} Electronics reviews show promising
6 directions. Using TF-IDF features (max 1,000 features), Multinomial Naive Bayes
7 achieved {nb_acc*100:.1f}% accuracy with MAE of {nb_mae:.2f}, while Logistic Regression
8 (nominal treatment) achieved {lr_acc*100:.1f}% accuracy with MAE of {lr_mae:.2f}.
9 Confusion matrix analysis reveals that {lr_adj:.0f}% of misclassifications occur
10 between adjacent ratings (particularly 4→5 stars), supporting our hypothesis that
11 ordinal treatment may improve performance. The class imbalance is evident—models
12 achieve {f1_5star*100:.0f}% F1-score for 5-star reviews but only {f1_3star*100:.0f}%
13 for 3-star reviews. These preliminary findings motivate our investigation into
14 whether ordinal methods can reduce MAE by better modeling rating structure while
15 addressing the adjacent-rating confusion problem."""
16
17 print("\n" + paragraph)

```

Initial experiments on 9,991 Electronics reviews show promising directions. Using TF-IDF features (max 1,000 features), Multinomial Naive Bayes achieved 60.5% accuracy with MAE of 0.74, while Logistic Regression (nominal treatment) achieved 62.6% accuracy with MAE of 0.62. Confusion matrix analysis reveals that 62% of misclassifications occur between adjacent ratings (particularly 4→5 stars), supporting our hypothesis that ordinal treatment may improve performance. The class imbalance is evident—models achieve 78% F1-score for 5-star reviews but only 6% for 3-star reviews. These preliminary findings motivate our investigation into whether ordinal methods can reduce MAE by better modeling rating structure while addressing the adjacent-rating confusion problem.